

Powering Votes?

The Effects of Energy Subsidies on Electoral Outcomes

Extended Abstract

Francisco Pizzi*

June 2026

1 Introduction

Energy subsidies are widespread in both developing and developed countries despite a broad consensus that they are a poor policy instrument (Black et al., 2023; World-Bank, 2025). They are fiscally costly and crowd out spending on better-targeted components of the social safety net; and when they subsidise fossil-based energy they reinforce the consumption of goods with large negative externalities (Burlig and Sudarshan, 2026). One reason these subsidies persist in spite of this consensus is political: removing or reducing them is salient, visible on every monthly bill, and routinely contested at the ballot box, so governments face political constraints when attempting to reform them. Yet the empirical magnitude of those political costs — how much electoral support an incumbent actually gains from granting an energy subsidy, or loses from withdrawing one — is still poorly understood.

Identifying a causal electoral effect of an energy subsidy faces two main challenges. First, subsidies are not randomly assigned: they are typically targeted to politically strategic regions or population groups, so a cross-sectional comparison between recipients and non-recipients would confound the effect of the subsidy with the unobserved political characteristics underlying both groups. Second, identification requires the source of treatment variation to operate at the level at which electoral outcomes are observed. Because ballot secrecy forces the use of aggregated vote tallies, random assignment at the individual level would not suffice: a credible design must deliver quasi-random variation in subsidy status across the units at which votes are counted, not merely across individuals.

*PhD candidate, Department of Economics, University of California, Davis. fmpizzi@ucdavis.edu.

This paper exploits an institutional feature of Argentina’s residential gas market that addresses both identification concerns. In 2021, midway through a legislative election year, the Argentine government extended an existing residential gas subsidy — historically restricted to the cold southern provinces — to a new set of cold regions in the centre and north of the country. Because this newly covered area is defined by a pre-existing *bio-climatic* map drawn years before the reform, the extension generated quasi-random variation in subsidy status across a number of new electoral circuits, unrelated to local political conditions. I exploit this boundary in a geographic regression discontinuity (RD) design, comparing the *change* in vote share in electoral circuits on either side of the boundary while controlling for the signed distance to it. The outcome is the difference between the share of valid votes received by the policy-associated party — the Kirchnerist/Peronist alliance — in the 2021 and 2017 general legislative elections. The unit of analysis is the electoral *circuit*, the finest administrative aggregation that is consistent across electoral years.

The main results point to a positive electoral return to the subsidy. On average, electoral circuits that received it show an increase in support for the policy-associated party roughly 5 percentage points larger than in circuits that did not. The sign of the estimate is consistently positive across specifications, although its magnitude varies substantially and statistical significance is mixed. Including census covariates — the share of households with unsatisfied basic needs, the unemployment rate, the share of households cooking with piped gas, and the share of highly educated households — leaves the point estimates largely unchanged, consistent with the boundary being uncorrelated with local economic conditions.

This paper contributes to the empirical literature on electoral cycles in the energy sector (Min and Golden, 2014; Mahadevan, 2024; Baskaran et al., 2015) by measuring the political impact of a highly salient subsidy granted in the middle of an electoral process. It also contributes to the literature on electoral rewards to incumbents (Manacorda et al., 2011), extending evidence established largely for cash transfers to a salient, in-kind energy subsidy, and to work on Argentine politics specifically (Ajzenman and Durante, 2023; Feierherd and Lucardi, 2023; Freille and Mazzalay, 2018).

The rest of the paper is organised as follows. Section 2 describes the institutional setting and the Cold Zone Law extension. Section 3 presents the data and sample construction. Section 4 lays out the empirical strategy, and Section 5 reports the results. Finally, Section 6 concludes.


2 The Cold Zone Law and its 2021 extension

Since 2003, residential gas consumers in a set of southern (cold) provinces have benefited from lower regulated prices than in the rest of the country. The subsidy is financed through a tax levied on all other residential gas consumers, and operates by compensating distributors that charge below-cost tariffs in the subsidised regions. The original rationale was geographic and climatic: in the southern part of the country, households face colder winters and therefore higher heating costs, which the

policy partially offsets. The boundary of the original subsidised area was based on a pre-existing bio-climatic classification of the territory. In 2021 — an electoral year for the national legislature — the government passed a law extending the subsidised area to additional regions classified as cold under the same bio-climatic map. The stated justification was to compensate historically cold consumers who had been excluded from the original 2003 scheme. The extension reached approximately three million new customers and reduced their gas bills by 30–50%.¹

Three features of the policy make it useful for empirical work. First, the spatial boundary of the new subsidised area was *not* drawn by the policymaker in 2021; it follows the pre-existing bio-climatic map (specifically, the addition of bio-climatic zones IIIa and IVa). Figure 1 shows an official government announcement publicising the extension, together with the bio-climatic map used to define its geographic scope.

Figure 1: The policy map



(a) Before/after of the Cold Zone area (official announcement).

(b) Bio-climatic map of Argentina. The 2021 extension incorporated zones IIIa and IVa (light orange).

Panel (a) shows the area covered by the subsidy before and after the 2021 extension, as published by the government. Panel (b) shows the bio-climatic map that defines the relevant climatic zones, drawn years before the 2021 change.

Second, energy subsidies have long been a contested issue in Argentina’s political-economic landscape (Cont et al., 2021; Navajas, 2015), with a clear ideological divide between parties on whether they should be reduced or maintained. This division is particularly relevant here because

¹Existing work has found that the extension did little to reduce energy poverty (Poggiere and Ibañez Martin, 2024) and had a weak redistributive effect, as is common with energy subsidies more generally (Bertín et al., 2025; Giuliano et al., 2020; Puig and Salinardi, 2015; Puig et al., 2025; Puig, 2023).

Kirchnerism/Peronism, the alliance that sponsored the 2021 extension, is consistently identified with the defence and expansion of energy subsidies, while the main opposition has historically argued for their reduction. The vote share of the policy-associated party is therefore a natural outcome to examine, since the subsidy is unambiguously associated with one side of the political spectrum.

Third, the rollout of the policy aligned closely with the electoral calendar. Figure 2 summarises the timing of the 2021 reform relative to the electoral calendar. The bill was introduced in Congress on May 14, 2021; it passed the lower chamber on June 10 and cleared the Senate on June 24. The new tariff schedule was published on August 5. The mandatory open primary elections (PASO) took place on September 12, and the general legislative elections on November 14. The new tariffs were therefore in force by the time the campaign for the general election was under way.

Figure 2: Schemathic Policy Timeline



The legislative election year of the extension lines up with the publication of the new tariff schedule and the two electoral dates.

3 Data and sample construction

Electoral data. The electoral data come from the datasets published by the Argentine *Cámara Nacional Electoral* (CNE) for every legislative and presidential election. For each election the dataset reports vote counts and registered voters at the most disaggregated level (the polling table, or *mesa*). The administrative hierarchy used in Argentine elections has six levels:

District → Provincial Section → Section → Circuit → Polling Place → Polling Table.

Polling tables are renumbered across elections and polling-place identifiers are inconsistent. The *circuit* is the finest level that is stable across years and can therefore be followed over time; it is the unit of analysis for the rest of the paper. I restrict attention to the *House of Representatives* ballot at the circuit level and construct the share of valid votes received by the policy-associated party. Defining the “treated party” is non-trivial because Argentine parties change names, alliances and composition across elections and across provinces. I follow the rule that the party associated with the policy is, in each election and each district, the list belonging to the Kirchnerist/Peronist

alliance. For the mandatory open primaries (PASO), all lists associated with the alliance count towards the party vote.

Geographic data. The circuit-level polygons are obtained from the *Dirección Nacional Electoral* (DINE-CNE)². The bio-climatic boundary is digitised from the climate map referenced in the law. Each circuit is assigned to one of three treatment categories according to whether its centroid falls inside the relevant region: circuits that always received the subsidy, circuits that never received the subsidy, and circuits that began to receive the subsidy after the reform. Figure 3 shows the spatial distribution of the three categories. The treated (new-subsidy) circuits form a band around the boundary between the original subsidised area and the cold extension defined by the bio-climatic map (ENARGAS, 2022). Table 1 summarises the number of provinces, sections, circuits, polling tables and registered voters in each treatment category for the reference election (Generales 2021).

Census data. Covariates come from the national census conducted in 2022, available at the census-tract level. They include the share of households cooking with piped gas, the share of highly educated households, the share of households with unsatisfied basic needs, and the unemployment rate. To probe the smoothness assumption underlying the RD and to include controls in robustness specifications, I assign the tract-level census covariates to circuits, dropping any census tract that spans more than one circuit.

4 Empirical strategy

The analysis sample is restricted to the *multi-treatment subset* of provinces — those that contain circuits in more than one treatment category — because provinces in which all circuits share the same status carry no within-province RD identifying variation. The always-subsidy circuits are dropped from the running-variable sample, so the comparison is between *New-subsidy* (treated) and *Never-subsidy* (control) circuits within the multi-treatment provinces.

The outcome is the change in the share of valid votes received by the policy-associated party between the 2017 and 2021 general legislative elections

$$\Delta y_i = y_{i,2021} - y_{i,2017}, \quad y_{i,e} = \frac{\text{Party votes}_{i,e}}{\text{Valid votes}_{i,e}} \in [0, 1],$$

where i indexes circuits and e indexes elections.

To remove cross-province variation in the change of baseline support for the policy-associated party, Δy_i is residualised on province fixed effects (estimated on the subset). Figure 4 compares the raw mean and interquartile range of Δy_i across provinces; Figure 6 restricts to the multi-treatment subset and shows the residualised version that the RD uses.

The baseline specification is a standard local-polynomial RD:

²Gaps in DINE coverage are filled with circuit polygons from the unofficial electoral repositories..

$$\Delta y_i = \alpha + \beta T_i + f(d_i) + T_i \cdot f(d_i) + \mu_{b(i)} + \varepsilon_i, \quad (1)$$

where $T_i = \mathbf{1}\{d_i \geq 0\}$ indicates the treated (new-subsidy) side, d_i is the signed distance in kilometres from the centroid of circuit i to the closest point on the new border, f is a local polynomial of order $p \in \{1, 2\}$, and $\mu_{b(i)}$ is an optional border-point fixed effect³. Several specifications of the model are run, varying bandwidth, order of polynomials and kernel weights. All standard errors are clustered at the section level.

Identifying assumption. The coefficient of interest, β , captures the discontinuous change in Δy_i at the new border. The identifying assumption is that, absent the extension, the change in vote share between 2017 and 2021 would not jump at $d = 0$, so that no discontinuity in Δy_i would arise. Under this assumption, any jump in Δy_i at the cutoff is attributable to the policy and identifies the average effect of being on the new-subsidy side. The assumption fails if demographics or unobservables evolved differentially across the border over 2017–2021, in which case the differential trend loads onto β . Adding the census covariates controls for differences in observable characteristics in levels, but not for their time trend, since the 2022 census provides essentially a single cross-section rather than a panel.

A further concern is that treatment is assigned at the section level, a coarser unit than the circuit at which the outcome is measured. If local governments respond to the policy, other section-level factors unrelated to the subsidy itself may change at the same border and be bundled with it, so that the estimated jump captures the subsidy together with any co-located institutional differences. This is the compound treatment problem familiar from geographic RD (Keele and Titiunik, 2015). I cluster the standard errors at the section level — the level at which the policy is assigned — to account for the resulting within-section correlation, and being able to observe what differs across the border helps argue that the subsidy, rather than some other section-level factor, drives the estimated jump.

Another concern affects the assignment of the treatment. Treatment is assigned by the location of the circuit, but the subsidy accrues to households that consume residential gas at a given address, so the design relies on people voting in the same place where they consume gas. If voter addresses are not kept up to date, some voters counted in a treated circuit do not actually receive the subsidy there, and the clean mapping from location to treatment breaks down. This is the usual distinction between a sharp and a fuzzy RD design: with stale addresses, being on the treated side raises the probability of being subsidised rather than guaranteeing it. The problem becomes more serious if the mismatch is systematically correlated with an omitted variable, since then it biases β rather than merely attenuating the estimate.

A related concern is anticipation and sorting. If people anticipated the 2021 extension, they

³The border is discretised into a small number of fixed points so that each circuit is anchored to its closest segment of the boundary (Keele and Titiunik, 2015; Cattaneo et al., 2026).

could have moved across the border, or re-registered their addresses, in order to land on the subsidised side. This undermines the as-good-as-random assignment of circuits near the border — the standard concern about manipulation of the running variable in an RD design. As before, the most damaging case is when this re-registration is systematic, that is, correlated with political preferences or other determinants of the vote, because then the circuits on either side of the cutoff are no longer comparable.

A further concern is the possibility of swinging alliances. If there are parties that change sides from one election to the other, then the comparison would be problematic. The use of the national deputies ballot, together with province and time fixed effects, alleviates this concern.

Finally, given the existing proposals to change the border again, there are two things to notice about the external validity of the paper. On the one hand, this was a period in which energy subsidies were considered less inefficient than in the current policy debate. Since 2023, parties have moved toward a more balanced fiscal approach, likely derived from their representation in the policy. This difference might matter for how people think about the policy and the rewards they give to politicians. On the other hand, during that period the energy bill in real terms was lower, so even if today energy subsidies might be considered less distortionary, their real weight on the household budget share is higher, contributing to a potentially higher effect on voter reactions.

5 Results

Figure 5 reports the baseline RD plots within a bandwidth of 150 km of the new border. Each panel plots the binned scatter of the change in the party’s vote share, Δy_i , against the signed distance of each circuit to the new border, overlaid with the local-polynomial fit and its 95% confidence band. The panels differ in the choice of the fitted polynomial, which is either linear or quadratic, and in the weighting kernel, which is either uniform or triangular. In every panel the fitted line jumps upward as it crosses the border into the new-subsidy side, though the size of this discontinuity is small relative to the dispersion of the underlying data.

Tables 2, 3, and 4 report the formal RD estimates. The three tables present different specifications for the estimation including different kernels, order of polynomials and border fixed effects. Table 2 is the baseline, estimated within a bandwidth of 150 km of the border. Table 3 adds the four census covariates — the share of households cooking with piped gas, the share of high-educated households, the share with unsatisfied basic needs, and the unemployment rate. Table 4 keeps these covariates but replaces the fixed bandwidth with the MSE-optimal two-sided bandwidth (Calonico et al., 2014).

Across the different specifications, the estimated effect at the border is positive. The magnitude varies across columns and across tables, and statistical significance is mixed — some specifications are significant at conventional levels and others are not — but the sign of the effect is consistent

throughout. The median coefficient across the models is five percentage points. Adding the census covariates leaves the point estimates broadly unchanged relative to the baseline, suggesting that the covariates do not absorb much of the residual variation around the cutoff.

The implication of the result is that the circuits that received the subsidy increased their support to the party associated with the policy by five percentage points more compared to the ones that did not receive the subsidy.

6 Conclusion


This paper asks how much electoral support an incumbent gains from granting a salient energy subsidy, a question that is hard to answer given the lack of randomness in subsidy allocation. This paper exploits the 2021 extension of Argentina’s *Ley de Zona Fría*, which enlarged a residential gas subsidy in the middle of a legislative election year along a boundary inherited from a pre-existing bio-climatic map. Because that boundary was drawn years before the policy and operates *across* electoral circuits rather than within them, it provides quasi-random variation in subsidy status aligned with the level at which votes are counted. Comparing the change in the policy-associated party’s vote share between the 2017 and 2021 general legislative elections across circuits on either side of the new border, I estimate that newly-subsidised circuits increased their support for the Kirchnerist/Peronist alliance by roughly five percentage points relative to circuits that remained excluded.

The estimated effect is positive for several specifications, with little adjustment when four census covariates are added. At the same time, the magnitude varies across bandwidths, kernels, and polynomial orders, and statistical significance is mixed, so the result is best read as suggestive evidence of a sizeable but imprecisely estimated electoral reward. Substantively, the finding speaks to the political-economy puzzle of why energy subsidies persist despite a broad consensus that they are fiscally costly, regressive, and environmentally distortive: if voters reward incumbents who extend them, governments face a concrete electoral incentive to maintain them, which helps rationalise their stickiness. The paper thus contributes a credible, geographically identified estimate of the political returns to energy subsidies.

Several caveats temper these conclusions and define the agenda going forward. Treatment is assigned at the section level while the outcome is observed at the circuit level, so the design carries the usual compound-treatment concern of geographic RD; identification also requires that no unobservable evolve differentially across the border, which the current specification can address in levels but not in trends, and that subsidised consumers actually vote where they consume gas—an assumption that would convert the design from sharp to fuzzy if voter addresses are stale or selectively updated. External validity is a further limitation: the estimate comes from a period in which energy subsidies were more politically normalised and real energy bills were lower than


today, so the contemporary electoral reward to such a policy could differ in either direction. Future work should sharpen inference by exploiting additional pre-policy elections to test for differential pre-trends, probe the residence-versus-voting margin directly, and assess robustness to alternative definitions of the policy-associated party.

Figure 3: Electoral circuits and treatment status




Circuits are coloured by treatment status: **green** for always-subsidy circuits, **blue** for new-subsidy circuits, and **red** for never-subsidy circuits. Own elaboration based on DINE and Enargas (2022).

Figure 4: Electoral results for party identified with the policy, by province



Note: Mean (diamonds) and $p25-p75$ (boxes) of $\Delta y_i = y_{i,2021} - y_{i,2017}$ by province. Province labels are coloured by treatment-status mix: black denotes a single status, blue denotes two statuses, red denotes three statuses.

Figure 5: Regression Discontinuity Plots



Four-panel RD on $\Delta y_i = y_{i,2021} - y_{i,2017}$ using the native `rdplot` output from the `rdrobust` package (Calonico, Cattaneo, Titiunik): bin-scatter plus local-polynomial fit and 95% confidence-interval ribbon, with the package's default styling. Panels span $p \in \{1, 2\}$ and uniform/triangular kernels. Bandwidth $h = 150$ km.


References

- Ajzenman, N. and Durante, R. (2023). Salience and accountability: School infrastructure and last-minute electoral punishment. *The Economic Journal*, 133(649):460–476.
- Baskaran, T., Min, B., and Uppal, Y. (2015). Election cycles and electricity provision: Evidence from a quasi-experiment with indian special elections. *Journal of Public Economics*, 126:64–73.
- Bertín, O., García, T., Pizzi, F., Porto, A., Puig, J., and Puig, J. (2025). Revisiting distributional effects of energy subsidies in argentina: O. bertín et al. *International Review of Economics*, 72(1):10.
- Black, M. S., Liu, A. A., Parry, I. W., and Vernon, N. (2023). *IMF fossil fuel subsidies data: 2023 update*. International Monetary Fund.
- Burlig, F. and Sudarshan, A. (2026). Reforming energy subsidies. *Handbook of Social Protection in the Developing World*, MIT Press, Forthcoming. Forthcoming.
- Calonico, S., Cattaneo, M. D., and Titiunik, R. (2014). Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica*, 82(6):2295–2326.
- Cattaneo, M. D., Titiunik, R., and Yu, R. R. (2026). Estimation and inference in boundary discontinuity designs. *arXiv e-prints*, pages arXiv–2505.
- Cont, W., Navajas, F. H., Pizzi, F., and Porto, A. (2021). *Precios y tarifas y política económica*. Facultad de Ciencias Económicas (UNLP).
- ENARGAS (2022). Ampliacion del régimen de zona fría: Ley 27.637, evolución de beneficiarios y facturas medias estimadas con niveles de ahorro. *Informe Técnico - Julio 2022*.
- Feierherd, G. and Lucardi, A. (2023). When the partisan becomes personal: mayoral incumbency effects in buenos aires, 1983–2019. *Journal of Elections, Public Opinion and Parties*, 33(4):684–704.
- Freille, S. and Mazzalay, V. (2018). Determinants of party and mayor reelection in local governments: An empirical examination for argentina during 1983-2011. *Revista de Economía y Estadística*, 56(1):59–87.
- Giuliano, F., Lugo, M. A., Masut, A., and Puig, J. (2020). Distributional effects of reducing energy subsidies: Evidence from recent policy reform in argentina. *Energy Economics*, 92:104980.
- Keele, L. J. and Titiunik, R. (2015). Geographic boundaries as regression discontinuities. *Political Analysis*, 23(1):127–155.

- Mahadevan, M. (2024). The price of power: Costs of political corruption in indian electricity. *American Economic Review*, 114(10):3314–3344.
- Manacorda, M., Miguel, E., and Vigorito, A. (2011). Government transfers and political support. *American Economic Journal: Applied Economics*, 3(3):1–28.
- Min, B. and Golden, M. (2014). Electoral cycles in electricity losses in india. *Energy Policy*, 65:619–625.
- Navajas, F. (2015). Subsidios a la energía, devaluación y precios. *documento de trabajo*, (122).
- Poggiese, M. and Ibañez Martin, M. (2024). Combining energy subsidies is not free: distributional effects and energy poverty. *LIX Reunión Anual de la AAEP*.
- Puig, J., Gasparini, L., and Puig, J. (2025). Measuring energy poverty in argentina: The relevance of regional disparities and policies. *Utilities Policy*, 97:102042.
- Puig, J. M. (2023). *Pobreza energética en Argentina*. PhD thesis, Universidad Nacional de La Plata.
- Puig, J. P. and Salinardi, L. H. A. (2015). Argentina y los subsidios a los servicios públicos: un estudio de incidencia distributiva. *Documentos de trabajo del cedlas*.
- World-Bank (2025). *State and Trends of Carbon Pricing 2025*. World Bank, Washington DC.

Annex

Figure 6: Electoral results for party identified with the policy residualized from province fixed effects, selected provinces



Note: Restricted to the multi-treatment subset of provinces from Figure 4, with Δy_i residualised on province fixed effects (grand mean added back to preserve the y -axis scale). Only within-province variation remains.

	Always subsidy	New subsidy	Never subsidy	Total
Provinces	9	10	16	23
Sections	82	160	237	479
Circuits	558	1,305	2,095	3,956
Tables	6,528	26,001	53,514	86,043
Registered voters	2,124,380	8,784,181	18,297,432	29,205,993

Table 1: Baseline counts by treatment status, Generales 2021. The figures show the size of each treatment group along the nested administrative hierarchy.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
RD est.	0.027	0.037*	0.062*	0.079**	0.026	0.031	0.060*	0.072**
SE	(0.020)	(0.022)	(0.034)	(0.032)	(0.021)	(0.021)	(0.034)	(0.029)
R^2	0.005	0.009	0.007	0.012	0.028	0.033	0.030	0.036
N_{left}	1320	1320	1320	1320	1320	1320	1320	1320
N_{right}	1056	1056	1056	1056	1056	1056	1056	1056
p	1	1	2	2	1	1	2	2
Kernel	U	T	U	T	U	T	U	T
h (km)	150	150	150	150	150	150	150	150
Border FE	N	N	N	N	Y	Y	Y	Y

Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$ (Conventional from `rdrobust`; SE clustered by sección, `gid.x`).

Kernel: U = uniform, T = triangular. Border FE: Y = yes, N = no. Outcome: Δ party share (2021 - 2017).

Table 2: RD specifications on $\Delta y_i = y_{i,2021} - y_{i,2017}$. Eight columns = {Border FE: N, Y} $\times$ { $p \in 1, 2$ } $\times$ {kernel: U, T}, bandwidth $h = 150$ km. Conventional point estimates and standard errors from `rdrobust`; border-FE columns pass border-id dummies through `covs`.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
RD est.	0.027	0.035	0.063*	0.074**	0.023	0.028	0.056*	0.069**
SE	(0.020)	(0.022)	(0.034)	(0.032)	(0.020)	(0.020)	(0.033)	(0.030)
R^2	0.014	0.029	0.017	0.032	0.042	0.062	0.045	0.065
N_{left}	1253	1253	1253	1253	1253	1253	1253	1253
N_{right}	1025	1025	1025	1025	1025	1025	1025	1025
p	1	1	2	2	1	1	2	2
Kernel	U	T	U	T	U	T	U	T
h (km)	150	150	150	150	150	150	150	150
Border FE	N	N	N	N	Y	Y	Y	Y

Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$ (Conventional from `rdrobust`; SE clustered by sección, `gid.x`).

Kernel: U = uniform, T = triangular. Border FE: Y = yes, N = no.

Outcome: Δ party share (2021 - 2017). All specs include census covariates via `covs`: `shr_gas`, `shr_high_educated_households`, `share_NBI`, `shr_unemployment`.

Table 3: RD specifications with census covariates on $\Delta y_i = y_{i,2021} - y_{i,2017}$. Same eight-column layout as Table 2; outcome residualised on province fixed effects on the working sample.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
RD est.	0.060*	0.057*	0.057	0.060	0.091***	0.087**	0.052	0.120**
SE	(0.032)	(0.030)	(0.039)	(0.039)	(0.033)	(0.040)	(0.053)	(0.059)
R^2	0.080	0.092	0.026	0.047	0.142	0.234	0.128	0.147
N_{left}	380	435	691	683	326	136	632	278
N_{right}	492	603	899	1094	180	354	351	927
p	1	1	2	2	1	1	2	2
Kernel	U	T	U	T	U	T	U	T
h_l/h_r (km)	52.4/67.4	56.5/80.6	79.6/123.6	78.2/169.1	47.6/43.0	32.8/56.7	73.8/56.3	42.0/131.5
Border FE	N	N	N	N	Y	Y	Y	Y

Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$ (Conventional from `rdrobust`; SE clustered by sección, `gid.x`).

Kernel: U = uniform, T = triangular. Border FE: Y = yes, N = no. Bandwidth: `bwselect=msetwo`.

Outcome: Δ party share (2021 - 2017). All specs include census covariates via `covs`: `shr_gas`, `shr_high.educated.households`, `share_NBI`, `shr_unemployment`.

Table 4: RD specifications with census covariates and MSE-two-sided optimal bandwidth (`bwselect=msetwo`) on $\Delta y_i = y_{i,2021} - y_{i,2017}$. Same eight-column layout as Table 3; outcome residualised on province fixed effects on the working sample.